

# LLM Inference Throughput Comparison

## NVIDIA A10 GPU (24GB)

